

Towards the Ultimate Universe Simulation

Introduction

Simulating the entire universe with perfect fidelity is an unprecedented scientific and engineering challenge. It demands modeling **all known physics** – from general relativity and the Standard Model of particle physics to quantum mechanics and thermodynamics – and potentially including **speculative theories** (string theory, loop quantum gravity, multiverse models) where justified by science. The vision is to create a digital cosmos starting at the **Big Bang**, evolving through cosmic structure formation and the emergence of life and intelligence, ultimately enabling **artificial life forms to arise from first principles**. Performance is a secondary concern to accuracy; however, practical constraints require clever trade-offs in scale and detail. This report assesses the current state of the `expressionneuroscienceinstitute/evolve` project (dev branch) and provides recommendations – backed by peer-reviewed science – for advancing it into the **most advanced universe simulation ever built**. Key topics include the integration of fundamental physics, bridging between scales, leveraging existing simulation frameworks, and best practices for multi-scale modeling. A comparison table of relevant open-source simulation engines is also provided to guide development decisions.

Current State of the EVOLVE Project

The EVOLVE repository's `dev` branch reveals a modular Rust-based codebase with multiple components already in place ¹. The architecture is split into dedicated crates for different concerns: a **physics engine** for fundamental particles and forces, a **universe simulation** module for celestial bodies and cosmic evolution, an **agent evolution** module for AI life and decision-making, networking code (for future distributed runs), diagnostics, plus a command-line interface (CLI) and a web-based visualization dashboard ². The project's README describes it as “the most advanced AI evolution simulation ever created – from fundamental particles to immortal intelligence” ³.

Implemented Features: Although still in development, EVOLVE has made tangible progress in several areas, emphasizing scientific realism:

- **Physics Engine:** The simulation prioritizes physics accuracy, using peer-reviewed physical laws and constants. The documentation notes use of *CODATA 2023 constants* and enforcement of conservation laws (energy, momentum, charge) ⁴. The code implements major equations of physics – Schrödinger's equation for quantum behavior, Maxwell's equations for electromagnetism, and even Einstein's field equations for gravity (at least in principle) ⁴. Recent commits show work on **nuclear and particle physics**: for example, adding realistic nuclear reaction models and cross-sections and refining the Klein-Nishina formula for photon-electron scattering ⁵ ⁶. There are constants for Standard Model particles (e.g. a Z-boson mass) and tests validating physics formulas against literature ⁷ ⁸. The team implemented a semi-empirical mass formula for nuclear binding energies and detailed stellar physics models (fusion, decay chains, neutron capture) based on peer-reviewed sources ⁹ ¹⁰. Thermodynamics is not neglected: entropy calculations and proper

thermodynamic evolution are being integrated ¹¹ to ensure, for instance, that the second law is respected in simulations of closed systems.

- **Cosmic Structure and Relativity:** The “universe_sim” module handles large-scale structure – star and galaxy formation, planetary systems, and environmental parameters. EVOLVE can generate celestial bodies and simulate gravitational dynamics (with plans to incorporate general relativity for high-mass or high-speed scenarios). In fact, the CLI already supports creating or deleting bodies (“**create-body**” with mass and type, “**delete-body**”) and adjusting physical constants on the fly ¹² ¹³, indicating a working gravitational framework. A rudimentary **ASCII universe map** feature is implemented to visualize density distributions of stars, gas, dark matter, etc., at different zoom levels ¹⁴. While it’s unclear if full Einstein field equation solvers are active, the design clearly aspires to relativistic gravity. The documentation promises that “*complex interactions that govern planetary formation*” are modeled with high fidelity ¹⁵.
- **Artificial Life and Agents:** A core goal is to watch AI-based lifeforms evolve within the physics simulation. The **agent_evolution** crate manages these digital organisms. Commits show progress on an AI decision-making engine with **sensory inputs, learning, and adaptive behavior** ¹⁶. The CLI allows inspecting agents or lineages in detail (“**inspect lineage <ID>**”) and even spawning new lineages via a “god mode” command (e.g. `spawn-lineage` with a given code hash on a target planet) ¹⁷. Each agent’s **decision history** and context can be tracked in the UI, and metrics like fitness, diversity, and innovation rates are recorded ¹⁸ ¹⁹. The simulation divides evolution into stages – from basic self-preserving agents up to “immortal intelligences” – and tracks milestones like the emergence of consciousness, technological advancement, and digital transcendence ²⁰ ²¹. This indicates that while the lowest-level biochemistry of life might not be fully simulated, there are abstracted evolutionary behaviors and possibly neural-network-like decision frameworks guiding the agents. Recent enhancements include more complex **interaction models** between agents (cooperation, competition) and improved learning algorithms ²² ²³.
- **Interface and Tooling:** EVOLVE can run in a **server mode** to allow real-time monitoring. The CLI (`universectl`) supports commands to start/pause the simulation, adjust speed, take state snapshots, etc. ²⁴ ²⁵. Many “**God-Mode**” commands exist for debugging or user intervention, such as performing “miracles” on planets (perhaps triggering custom events) ²⁶ or warping time by a factor ²⁷. The web dashboard (`viz_web` crate) connects via WebSockets to display simulation state graphically in real time. Recent development focused on making the WebSocket connection robust and improving the dashboard’s UI ²⁸ ²⁹. The dashboard includes multiple visualization modes and analytical panels: users can **pan and zoom** through the simulated universe, toggle layers (e.g. viewing just dark matter or radiation fields), and click on entities to get detailed stats ³⁰ ³¹. There are real-time readouts for global metrics (universe age, average temperature, energy density) and for evolution stats (population counts, number of civilizations, etc.) ³². A performance monitor shows FPS, CPU/memory usage, etc., reflecting the work done on a diagnostics module ³³.

Scientific Rigor: Impressively, the EVOLVE project emphasizes *scientific accuracy* throughout its documentation. It explicitly states that **peer-reviewed equations** underlie the simulation and that known physics is followed as closely as possible ⁴. For example, quantum mechanical behavior is intended to be proper wavefunction evolution with measurement mechanics accounted for, and the thermodynamic behavior ensures entropy increases appropriately ⁴. The team has built in **validation tests** – one commit mentions “*peer-reviewed cross-section validation*” for the physics engine ⁶, indicating that simulation

outputs (like reaction rates or scattering outcomes) are being checked against established scientific results. Evolutionary dynamics are likewise checked: the project references Hardy-Weinberg equilibrium in population genetics and realistic selection pressures to ensure that digital evolution isn't arbitrary but grounded in real evolutionary theory ³⁴ .

In summary, the current `dev` branch of EVOLVE provides a strong foundation with modules for physics, cosmic evolution, and AI agents. Many features are prototyped (or partially implemented): from nuclear fusion in stars to AI decision-making loops. The overall framework is in place to simulate a full universe timeline – divided into eras such as Particle Soup, Star Formation, Planetary Age, Biogenesis, and so on up to a far-future AI civilization epoch ³⁵ . The next steps will involve fleshing out each layer of physics and bridging them together seamlessly. Below, we outline recommendations and resources to guide EVOLVE's development into a true **first-principles universe simulator**.

Incorporating All Known Physics

A central requirement is that the simulation should include **all fundamental physics** that governs our universe, across all scales. This is extraordinarily ambitious: modern science spans a vast range of length and time scales, from Planck-scale phenomena ($\sim 10^{-35}$ m, 10^{-43} s) to cosmic scales ($\sim 10^{26}$ m, 10^{17} s). No single simulation can brute-force every detail at once due to computational limits (more on that in a later section). The key is to integrate each domain of physics in a **consistent, multi-scale fashion**. Below we discuss the major physics areas and how they can be modeled or approximated, with pointers to methodologies and references.

General Relativity and Gravitation

At cosmic and astrophysical scales, **gravity** (described by Einstein's general relativity) is the dominant force shaping structure. Any universe simulation must account for relativistic gravity to model phenomena like the expansion of the universe, black holes, gravitational lensing, or even just accurate orbital dynamics for massive bodies.

Approach: In practice, full solution of Einstein's field equations everywhere is computationally infeasible for a large, dynamic universe simulation. Instead, use **hybrid methods**: adopt Newtonian gravity or simplified relativistic models where appropriate, and reserve full general relativity for regimes that demand it (e.g. near supermassive black holes or during the very early universe at extreme densities). Many **cosmological simulation codes** successfully model structure formation using Newtonian gravity plus a relativistic background expansion. For example, the open-source code *GADGET* uses an N-body approach with a particle-mesh or tree algorithm to compute gravitational forces among millions to billions of particles, and has been used in landmark simulations of cosmic structure ³⁶ . *GADGET* (and successors like **GADGET-4** ³⁷) assume Newtonian gravity with periodic boundary conditions, which is accurate on small scales, while the cosmological expansion is handled by scaling positions according to Friedmann's equations (from GR). EVOLVE can employ a similar strategy for the large-scale universe: use a **particle-based gravity solver** (tree or fast multipole method) for efficiency, and incorporate cosmological parameters (Hubble expansion rate, dark energy equation of state) to simulate an expanding universe if desired.

For strong-field scenarios (neutron stars, black holes, etc.), consider integrating a specialized **relativistic solver** or using results from it. The **Einstein Toolkit** is an open-source framework for relativistic astrophysics simulations ³⁸ . It provides C++/Fortran code to solve Einstein's equations on a 3D grid (using

finite-difference or spectral methods) and has been used to simulate black hole mergers and gravitational waves. While running Einstein Toolkit concurrently with the entire universe simulation may not be feasible, EVOLVE could interface with it or incorporate a simplified version for certain regions. For instance, a black hole could be represented by an analytic metric solution (like Schwarzschild or Kerr metric) for its immediate vicinity, or by precomputed lookup tables for its effects on nearby matter and light.

It's also important to include **relativistic corrections** for fast-moving particles or spacecraft (if any). Special relativity (Lorentz time dilation, mass-energy equivalence) can be implemented in the physics engine's particle dynamics. The EVOLVE docs already list Einstein's equations among its key equations ⁴, so the intent is there. Ensuring that (at least in aggregate) energy and momentum are conserved and that the speed of light is not violated will keep the simulation physically reasonable. In summary, use Newtonian gravity for bulk N-body calculations (valid when velocities $\ll c$ and fields are weak), and introduce relativistic modules for **high-gravity or high-velocity** conditions.

Resources:

- *Einstein Toolkit* – community-driven code for general relativity, which includes modules for hydrodynamics and magnetohydrodynamics under strong gravity ³⁸. This could guide how to incorporate GR in a modular way.
- *Geodesic solvers* – algorithms to propagate particles or photons through a curved spacetime metric. These can be used for ray-tracing light paths (for visualization of gravitational lensing or black hole shadows) if needed.
- *Post-Newtonian formulas* – for scenarios like orbital dynamics around a spinning black hole, closed-form solutions or series expansions from general relativity (e.g. post-Newtonian corrections to gravity) can be implemented to get much of the effect without full field simulation.

Quantum Mechanics and the Standard Model

At the opposite end of the scale, a *faithful* universe simulation should capture quantum physics: the behavior of subatomic particles, atomic structure, and particle interactions per the Standard Model. This includes quantum field theory (QFT) aspects for particle creation/annihilation, forces like the strong and weak nuclear force, and quantum statistics. Directly simulating all quantum degrees of freedom is utterly infeasible – even a single atom's electron cloud typically requires solving Schrödinger's equation in a high-dimensional space. However, certain approaches can incorporate quantum effects effectively:

- **Effective/Statistical Modeling:** For many purposes, detailed quantum dynamics can be replaced with statistically correct models. For example, rather than simulate each photon and electron's wavefunction during Compton scattering, one can use the **Klein-Nishina formula** to probabilistically determine scattering angles and energy transfer ⁵. EVOLVE already does this for certain processes. Similarly, radioactive decay can be simulated by random draws from an exponential decay law with known half-lives, producing appropriate daughter particles.
- **Quantum Field Events:** The Standard Model processes (particle collisions, decays, fusion reactions, etc.) can be implemented via event generators. **Geant4**, a well-established open toolkit from CERN, is a prime example: it provides libraries of particles and interactions allowing simulation of how particles propagate and interact with matter ³⁹. Geant4 includes high-energy processes (like electron-positron annihilation, quark hadronization, cosmic ray cascades) as well as nuclear reactions and low-energy interactions, all based on experimental cross-sections and theoretical

models ³⁹. Rather than reinventing the wheel, EVOLVE could integrate a simplified version of Geant4 or use its data tables for particle interactions. For instance, if the simulation needs to simulate cosmic ray particles hitting a planet's atmosphere, Geant4's physics lists could be employed to generate the shower of secondary particles and energy deposition in a scientifically accurate way. (Geant4 is C++ based, but bindings or a custom Rust port could be explored if performance allows.)

- **Quantum Chemistry:** Quantum mechanics is crucial for atomic and molecular behavior, which underpins chemistry and hence biology. Simulating quantum chemistry from first principles (e.g. solving Schrödinger's equation for every electron in a cell's worth of atoms) is beyond reach. Instead, **approximate methods** from computational chemistry can be used. One approach is to implement a *simple quantum chemistry engine* for basic reactions and bonds – for example, a parameterized molecular dynamics model or a reactive force field. The EVOLVE physics engine could use classical potentials for molecular interactions (perhaps switching to a quantum treatment for very simple systems if needed). There are open-source quantum chemistry packages (e.g. **Quantum Espresso** for density functional theory, though that's heavy) that likely won't run inside the simulation due to their computational intensity. A more feasible route is **pre-compute and tabulate** important chemical reactions. For instance, the rates of prebiotic chemical reactions (like formation of amino acids from precursors) can be included based on lab data or quantum chemistry calculations done offline ⁴⁰. During simulation, whenever conditions for a reaction are met (molecules present, temperature/pressure suitable), the event can occur stochastically with the correct probability and outcome.
- **Quantum Statistics:** Ensure that macroscale properties emerging from quantum physics are captured. This includes using proper statistics for fermions vs bosons (Fermi-Dirac and Bose-Einstein distributions) in relevant regimes. For example, if simulating neutron star matter or early-universe primordial nucleosynthesis, quantum degeneracy pressure and nuclear reaction rates require quantum statistical input. EVOLVE can incorporate these via formulas or look-up tables from nuclear physics. The mention of using the semi-empirical mass formula ¹⁵ suggests nuclear binding energies and thus nucleosynthesis can be computed – indeed, a realistic stellar evolution model needs to follow element creation through fusion and neutron capture, which the commits indicate has been added ⁹. Those nuclear reaction networks are built on quantum physics results but can be simulated as coupled rate equations for isotopes.

In summary, include quantum effects **implicitly**. Only simulate wavefunctions or quantum superpositions explicitly if absolutely necessary (perhaps for a quantum computing experiment within the simulation?). Otherwise, rely on well-tested models and libraries: cross-sections, reaction rates, decay laws, etc., from the vast body of particle and nuclear physics research. By doing so, the simulation remains *faithful* to known physics without incurring the impossible computational cost of direct quantum simulation.

Resources:

- *Geant4* – as noted, an extensive toolkit for particle interactions ³⁹. Its databases of physical processes (from low-energy neutron scattering to high-energy collider physics) could hugely benefit EVOLVE's realism in particle physics.
- *Particle data tables* – the Particle Data Group (PDG) publishes authoritative tables of particle properties and interaction probabilities. These should be used for masses, lifetimes, branching ratios of decays, etc., ensuring the simulation's particle physics aligns with experimental reality ⁷.
- *Astrophysical reaction networks* – codes like *STARLIB* or *REACLIB* provide libraries of thermonuclear reaction

rates (e.g. for hydrogen fusion, helium capture, supernova nucleosynthesis). These can plug into the stellar evolution aspect of EVOLVE. For example, EVOLVE could simulate a star by treating concentric shells and applying these reaction rates to evolve composition and energy output over time (if not simulating every atom individually).

- *Quantum random number generators* – since quantum processes are stochastic, using high-quality randomness is important. The simulation might consider using a robust entropy source or even conceptually a virtual “quantum oracle” for randomness to mimic true quantum indeterminacy when choosing outcomes of decay or measurement.

Thermodynamics and Statistical Physics

Thermodynamics provides the link between microscale physics and macroscale behavior. Any large simulation will need to keep track of energy flows, heat, entropy, and phase transitions. EVOLVE has already identified **entropy calculation** as a task ¹¹, which is crucial for verifying the second law and realistic behavior (e.g. no free energy violations). Key considerations:

- **Energy Conservation:** Ensure that when particles interact or when fields do work on matter, energy is conserved (to the precision of the model). For example, if matter annihilates into radiation, the simulation should create photons carrying the appropriate energy. The documentation emphasizes preservation of energy and other conserved quantities ⁴. This may involve choosing sufficiently small time-steps or using symplectic integrators for dynamics to minimize numerical drift in energy.
- **Heat and Dissipation:** Many processes (friction, inelastic collisions, chemical reactions) convert organized energy into heat. Representing heat at a fundamental level means simulating the random motions of particles. However, at higher levels one might treat materials as having temperature and specific heat rather than tracking each phonon. Best practice is to allow **multi-scale thermal modeling**: e.g., treat a solid planet’s interior with bulk properties (thermal conductivity, heat capacity) rather than simulating every molecule, but derive those properties from the microscopic physics. The simulation could use finite-difference or finite-element methods to model heat diffusion in continuum regions. If the scale allows, incorporate **Boltzmann transport** equations or **Monte Carlo thermal transport** for photons (especially in stars, where radiative diffusion is key to energy transfer).
- **Entropy and Arrow of Time:** To simulate **thermodynamics** correctly, one must introduce irreversible processes or randomness. A purely deterministic Newtonian simulation of classical particles conserves entropy (Liouville’s theorem) unless coarse-grained. In EVOLVE, irreversibility will enter through many channels: inelastic collisions (modeled by generating random thermal motions), quantum measurements (which effectively increase entropy when entanglement is curtailed), or user-defined coarse-graining (like “low-memory mode” that may drop small-scale details and thus mimic entropy increase). The team can explicitly track entropy by summing contributions from subsystems (using $S = k_B \ln \Omega$ or via information entropy if treating the simulation’s state as information). They have added entropy calculations to the physics engine ¹¹, likely to monitor that overall entropy of a closed system doesn’t decrease spontaneously. For accuracy, it will help to validate that in test scenarios (like gas in a box expanding, or two objects reaching thermal equilibrium) the entropy change matches theoretical expectations.

- **Phase Transitions:** The simulation should capture important phase changes (e.g. gas \rightleftharpoons liquid \rightleftharpoons solid, ionization, plasma formation, etc.) since these affect planet formation, star interiors, and living environments. Rather than modeling every molecular bond, one can use **equations of state** from physics. For instance, use well-known models like the ideal gas law (and its extensions for high pressure), or the **Van der Waals equation** for real gases, or tabulated free energy curves for water's phase diagram. This way, if a region of the simulation hits certain pressure-temperature conditions, the code can decide that water condenses, releasing latent heat, etc. Similarly, in stellar evolution, incorporate phase transitions like hydrogen ignition (when core temperature crosses ~ 10 million K and fusion rapidly increases). Rigorously, these should emerge from microphysics, but practically they can be triggered via condition checks tied to known physics.

In short, thermodynamics in the simulation will emerge from the myriad interactions of particles (if simulated at micro level) but must also be tracked at macro level. Embedding the **laws of thermodynamics** (energy conservation, entropy increase) as guiding constraints will ensure no violations of physics even when approximations are made.

Resources:

- *Statistical physics references* – texts or data for thermodynamic properties of materials and mixtures across conditions (for example, the NIST database for material properties, or the OPAL opacity tables for stellar matter). Having these on hand helps to calibrate the simulation's macro behavior.
- *Heat transfer codes* – if needed, libraries for solving heat equations or handling radiative transfer (there are open codes used in astrophysics for radiative transfer which might be adapted).
- *Verification cases* – test the simulation on known thermodynamic scenarios: e.g., the Carnot cycle (should see no net entropy decrease), cooling of a hot object in a cold environment (compare with Newton's law of cooling), etc., to build confidence that the thermodynamics are correct.

Electromagnetism and Plasma Physics

Electromagnetic forces are crucial at atomic and molecular scales and also govern plasma behavior in stars and interstellar space. Maxwell's equations should be part of the simulation's core equations ⁴. A full electromagnetic solver (like solving Maxwell's PDEs on a grid) might be necessary for certain scenarios (e.g. simulating how a stellar magnetic field evolves, or how a laser by an advanced civilization might propagate). However, a grid-based EM simulation at universe scale is untenable. Instead, handle EM in a context-dependent manner:

- At particle level, use **Coulomb forces** for charged particles (e.g. protons, electrons) in close proximity, and include **magnetic Lorentz forces** for charges in motion. This could be done by pairwise interactions (which is expensive $O(n^2)$) or by methods like Particle-In-Cell (PIC) if simulating plasmas. PIC codes (like the open-source *OSIRIS* or *XOOPIC*) treat electromagnetic fields on a grid but particles as discrete, which might be integrated on sub-regions (say, to simulate a local plasma phenomenon).
- For larger continuous media, treat electromagnetic effects via **bulk properties**: electrical conductivity, magnetization, etc. For example, a planet's magnetic field could be approximated by a dipole field if it has a rotating iron core (the simulation could assign a dipole moment to the planet and simulate magnetospheric effects analytically rather than computing from first principles each charged particle in the core). In stars, include the effects of magnetic fields in a parameterized way

(perhaps enhancing convection or trapping charged particles). These are higher-level modeling decisions that can be guided by known astrophysical observations.

- If simulating technology (since AI civilizations might build devices), then electromagnetic field simulation might be needed for circuits or antennas. Rather than finite-difference time-domain solving Maxwell's equations everywhere, one could incorporate **circuit simulators** or simplified electronic models at the device level.

Given EVOLVE's scope, a reasonable approach is **implement Maxwell's equations but apply adaptively**. When simulating at atomic scale, the fields of individual particles can be computed by Coulomb's law. When simulating at continuum scale, solve Maxwell's equations on a coarse grid or via potential fields for regions of interest (with time-steps constrained by the speed of light for stability).

Resource: The **Athena++** code (open-source) is a versatile astrophysical magnetohydrodynamics code that includes **MHD (magneto-hydrodynamics)** with adaptive mesh refinement ⁴¹. It can simulate plasmas (e.g. accretion disks, stellar winds) using fluid equations plus Maxwell's equations. Studying Athena++ or similar codes (like **PLUTO** or **RAMSES** with MHD) could inform how to incorporate electromagnetic fields efficiently in a large simulation. For example, instead of tracking every photon, treat light as a field when wavelengths are small relative to system size, and switch to a ray or particle representation when needed (for instance, for high-energy photons or individual cosmic rays).

Unifying and Bridging Scales

One of the biggest challenges is bridging the **quantum-to-cosmic** gap. There is currently no complete "theory of everything" that unifies gravity with quantum mechanics, so even conceptually the simulation will have to make discrete choices about where one theory gives way to another. Some guidelines for a cohesive approach:

- Use **appropriate effective theories at each scale**, and ensure they feed into each other at interfaces. For example, use quantum mechanics to derive atomic properties (like bond strengths, emission spectra) and feed those into a classical chemistry simulation for molecules. Use the outcomes of stellar nuclear burning (from nuclear physics calculations) to feed a stellar structure model that produces luminosity and feedback into the galaxy. This way, each layer of the simulation is informed by the deeper layer's rigorous results without trying to compute everything bottom-up all the time.
- Adopt **cross-scale checkpoints**: At certain milestones, one can re-evaluate if lower-scale simulation needs to be refined. For instance, during "Biogenesis" (the era where life emerges), one might need a more detailed chemical simulation on habitable planets. EVOLVE could detect when conditions on a planet are suitable for life (temperature, presence of water and organics) and then *zoom in* on that planet with a more fine-grained simulation module (perhaps a Monte Carlo simulation of prebiotic chemistry, or an artificial life simulation at the cellular level). This is akin to **adaptive resolution**: not all parts of the universe need the highest detail at all times – only where interesting complexity is arising. This strategy aligns with the project's mention of **Level-of-Detail (LOD) adaptive resolution** ⁴², an important performance optimization. By dynamically adjusting the granularity of simulation (e.g. merging particles into a fluid description when detailed tracking isn't necessary, or splitting a fluid into particles when detail is needed), one can maintain fidelity where it matters most.

- Ensure **consistency conditions** between scales. For example, if you treat a planet’s atmosphere as a bulk fluid, but occasionally inject individual molecules for a detailed chemistry event, make sure that on average the bulk properties (pressure, temperature) match what the molecules are doing. This may require periodic synchronization between micro and macro models. Techniques like **coarse-grained molecular dynamics** (where groups of atoms are treated as single particles) might inspire how to smoothly interpolate between a fully atomic simulation and a continuum one.
- Consider the use of a **universal unit system** (likely natural units or SI units) across the simulation to avoid unit mismatches. The simulation should maintain a consistent set of fundamental constants (speed of light c , Planck’s constant, gravitational constant G , Boltzmann’s constant, etc.) so that when switching models, they’re all calibrated to the same reality. EVOLVE uses up-to-date CODATA values for constants ⁴, which is good practice.

In essence, treat the simulation as layered. Fundamental laws operate at the bottom (quantum fields, fundamental particles), emergent laws operate at higher levels (e.g. chemistry, thermodynamics, elasticity, fluid dynamics), and even higher, emergent behaviors like life and intelligence come into play. By **modularizing** these layers (as EVOLVE’s architecture already does with separate crates ¹), developers can work on each aspect with the appropriate scientific methods and then focus on the interfaces between modules.

Incorporating Speculative Theories

The brief calls for integrating “leading speculative theories (e.g. string theory, loop quantum gravity, causal dynamical triangulations, multiverse hypotheses) where scientifically defensible.” While these theories are unproven, including them could future-proof the simulation or allow experiments to see their potential effects. Here we discuss a few such theories and how one might include them in a simulation context:

- **String Theory:** In string theory, particles are one-dimensional strings, and extra spatial dimensions may exist (often 10 or 11 dimensions in total). Directly simulating string dynamics is far beyond current computing (even individual string vibrations are complex), and the theory lacks unique experimental predictions so far. However, one could incorporate **indirect consequences**: for example, certain string theory variants imply the existence of supersymmetry (every particle has a superpartner). If desired, the simulation could include supersymmetric particles in the physics engine (photinos, gluinos, etc.) and allow reactions that produce them, toggling this on as an option. Another possible inclusion is **additional forces or variations in constants**: some string models lead to very slow changes in fundamental constants or new force-carriers. If scientifically justified, EVOLVE could let the user enable an option where, say, the gravitational constant evolves slightly over cosmic time (per some scalar field theory), or there is a very weak fifth force acting on dark matter, etc., derived from beyond-Standard-Model physics. These would be speculative, so they should be turned off by default, but the framework can accommodate them modularly. The simulation’s architecture of modifiable constants ¹³ and its *god-mode* “set-constant” command ¹² suggest it’s relatively straightforward to allow alternate values or extra constants for experimental runs.
- **Loop Quantum Gravity (LQG):** LQG proposes that space is quantized into tiny “loops” or spin networks, leading to a smallest length scale on the order of Planck length ($\sim 1.6 \times 10^{-35}$ m). One way to incorporate this is to discretize spacetime at Planck-scale intervals. However, the number of

Planck cubes in even a single cubic meter is astronomically huge (on the order of 10^{105} per m^3), so a full grid is impossible. Instead, one could use LQG results in a statistical sense. LQG and related approaches like **Causal Dynamical Triangulations (CDT)** have had intriguing results: CDT simulations of quantum gravity in 4D showed that spacetime could spontaneously have a *fractal structure*, behaving effectively 2-dimensional at Planck scales, while emerging as 4D at large scales ⁴³. To integrate such findings, EVOLVE might allow an alternate gravity module where near the Planck scale, the behavior of gravity (or the dispersion of particles) changes according to these theories. For instance, one could modify the dispersion relations of particles at extremely high energies to reflect a granularity of space (some quantum gravity theories predict a slight energy-dependent speed of light, which could be simulated as a tiny change in photon speed at TeV energies – though this is speculative and not observed yet). Another possible feature is including a **minimum black hole size** (since LQG predicts black hole area quantization), meaning the simulation would not allow singularities below a certain mass/size. Again, these are experimental toggles that can be built in without being default. They would make the simulation a testbed for quantum gravity ideas – e.g., seeing if an LQG-inspired rule affects early-universe evolution or black hole evaporation outcomes.

- **Causal Dynamical Triangulations (CDT):** CDT is a concrete algorithmic approach to quantum gravity where spacetime is built from simple building blocks (simplices) glued in a way that preserves causality. One cannot simulate an entire universe at the atom level, let alone at the Planck level with CDT, but perhaps a *micro-simulation* of a patch of spacetime could be embedded. EVOLVE could have a mode to simulate a tiny region with a full CDT algorithm, to see emergent geometry. This might not influence the large-scale sim unless one wants to replace the classical continuum with a CDT-based dynamic lattice at all scales (which is computationally prohibitive for now). A more feasible inclusion is to use CDT insights to handle Planck-scale cutoffs. For example, as mentioned, CDT suggests spacetime is effectively 2D at the Planck scale ⁴³, which implies degrees of freedom “freeze out”. EVOLVE might mimic that by not attempting to represent spatial structure below some granularity (like it won't track distinct particles if they are within distances smaller than some epsilon, effectively treating them as one unit). This way, the simulation acknowledges a fundamental scale and could avoid singularities (like infinite density) by always “capping” density at roughly one particle per Planck volume or something similar.
- **Multiverse Scenarios:** The concept of a multiverse – multiple parallel universes with varying physical constants or initial conditions – is even more speculative. However, EVOLVE could allow a **multi-universe simulation** in two senses: (1) Simulate multiple independent universes in parallel (useful for ensemble experiments or comparing different physics). This is straightforward technically (just run several instances or threads with different parameter sets). (2) Simulate a branching quantum multiverse (like the Many-Worlds interpretation). The latter would mean that every quantum event can spawn separate outcomes. It's practically impossible to explicitly branch the entire simulation's state for every quantum measurement (the combinatorial explosion would be enormous). Instead, one could simulate **one branch at a time** but allow an observer within the simulation to perceive quantum uncertainty. Implementing Many-Worlds explicitly is not feasible; instead, one might simulate quantum probabilities and if needed, explore alternative outcomes by resetting to a checkpoint and randomizing different outcomes (i.e. effectively exploring branches sequentially rather than simultaneously). This is more of an analysis tool than a core simulation feature. Given performance priorities, multiverse features might be low priority, but EVOLVE's design with a flexible physics engine could be extended to different “universe settings” (for instance, you

could re-run the simulation with a different fundamental constant to see how things change – a kind of virtual multiverse experiment).

Summary: Integrating speculative physics should be done **modularly and optionally**. The default simulation should reflect known physics (to ensure it matches what we observe in our universe). Modules for speculative theories can be toggled for exploratory purposes. This means clearly isolating code that implements unverified physics. For example, have a compile-time or run-time flag for “Quantum Gravity Mode” that, when on, replaces the classical gravity calculations at very small scales with a chosen model (LQG, CDT, etc.). All such additions should be documented and sourced from actual theory papers to maintain scientific rigor. If done carefully, EVOLVE could become a powerful tool for researchers to test the implications of theories like string theory or loop quantum gravity on a simulated universe – something currently mostly thought about in theory or limited toy simulations.

Trade-offs: Scale, Fidelity, and Computational Cost

Absolute “perfect fidelity” down to the Planck scale for an entire universe is not achievable with any conceivable computer. The informational content of such a simulation is astronomically high. For perspective, physicist Seth Lloyd estimated that to simulate the entire observable universe **at the quantum level** would require on the order of 10^{90} bits of information and 10^{120} operations – which is essentially the maximum that the universe itself could have performed in its lifetime ⁴⁴ ⁴⁵. In other words, you’d need a computer as large as the universe running for billions of years to fully simulate it in detail. Therefore, EVOLVE must make **strategic trade-offs** in modeling fidelity versus computational load. Below we highlight key trade-offs and strategies to handle them:

- **Discrete vs Continuous Representation:** Representing matter as discrete particles everywhere (down to atoms or below) is extremely expensive. For example, an Earth-sized planet has $\sim 10^{50}$ atoms. Tracking each as an entity is impossible. Thus, use **continuous or collective representations** when appropriate. Treat solids and liquids as continuum materials (with fields for density, temperature, etc.), not as billions of atoms. Use particles only for what needs discrete resolution (perhaps gas in space can be “superparticles” each representing a clump of gas, as done in cosmological N-body simulations ³⁶). This drastically reduces the number of entities. Modern cosmological simulations represent the entire observable universe with on the order of 10^{10} particles for dark matter and gas – a huge number, but far less than the actual particle count in the universe, relying on each simulation particle to represent a huge collection of real particles. EVOLVE can employ similar **coarse-graining**. The trade-off is losing fine detail, but gaining tractability. The *Level-of-Detail (LOD)* approach mentioned in EVOLVE’s plan ⁴² encapsulates this: dynamically adjust whether a region is simulated as individual entities or as a smoothed continuum.
- **Spatial and Temporal Resolution:** The simulation must choose appropriate time-steps and spatial cell sizes. To cover **13.8 billion years** of evolution, the time-step can’t be too fine or the simulation would never finish. EVOLVE’s CLI has an option `--tick-span <YEARS>` ⁴⁶, implying the simulation advances in increments of years (or configurable span). For early epochs like the Big Bang, a year is too large – one might need sub-second resolution to simulate primordial nucleosynthesis (~first 3 minutes of universe) or the recombination era (hundreds of thousands of years after Big Bang) with any fidelity. A possible solution is **variable time-stepping**: use small ticks when events are rapid or conditions changing fast, and larger ticks in eras where evolution is slow (e.g., during intergalactic cruising between star formation peaks). This is analogous to what is done

in N-body codes that adapt the time-step based on forces or in stiff chemical reaction solvers that adapt to reaction rates.

Spatially, adaptivity is key: use a **dynamic grid or tree refinement** for areas of interest. For example, if a star is forming, concentrate resolution (particles or grid cells) in that region to capture the collapse, while keeping far-away interstellar space at coarse resolution. Adaptive Mesh Refinement (AMR) techniques, as employed by codes like **Enzo** or **RAMSES**, could be incorporated ⁴⁷ ⁴⁸. They start with a coarse grid for the whole simulation volume and refine subdividing cells where needed (high density or strong gradients). EVOLVE might integrate a simple AMR for its continuum parts, or an oct-tree for particle clustering (like Barnes-Hut tree codes do). The **trade-off** is additional code complexity and ensuring consistency across levels of refinement, but the benefit is huge: orders-of-magnitude savings in computation by not over-resolving empty or uniform areas.

- **Parallelism and Hardware Utilization:** To maximize performance (even though it's secondary to accuracy, it is still necessary for a runnable simulation), EVOLVE should exploit parallel computing. Rust's ecosystem and the mention of **multi-threading and ECS (Entity-Component-System) architecture** ⁴⁹ suggest the project is already aiming for high performance. An ECS design is great for organizing simulation data and making parallel updates efficient. The team should ensure thread-safe handling of shared data (likely using Rust's concurrency features to avoid data races). For large scale, eventually distributing the simulation across a **cluster** or using GPU acceleration might be needed, especially for the physics engine (e.g., collision computations, field solves). EVOLVE's networking crate hints at future **distributed simulation** support ⁵⁰, perhaps to run parts of the sim on different machines (e.g., each machine simulating one galaxy or one region of space). This is a complex but worthwhile direction – many advanced simulations (climate models, cosmology, etc.) run on supercomputers with MPI (message-passing interface) to communicate. EVOLVE could use a high-level approach (since it's in Rust, something like Rayon for threads and maybe an MPI-binding for multi-node).

A caution: distributing a **universe simulation** must handle the **speed-of-light limit** for information – one cannot update distant regions faster than light travel in a real physical sense. However, many simulations simply ignore this for practicality (updating globally each tick). If extremely high fidelity is desired, one could impose a light-speed constraint in the code (e.g., don't allow an event in one region to influence another region in less time than light would take). This could be approximated by only syncing neighboring regions frequently and distant regions less frequently (introducing a concept of signal travel). This is advanced and usually not done explicitly in simulations (except those focusing on relativity or computational causality). Still, it's something to bear in mind as a fidelity concern.

- **Data Management:** A simulation of this scope generates an enormous amount of data. Storing every particle's history would be prohibitive. EVOLVE's plan to compress time-series data ⁵¹ is smart. Developers should implement **checkpointing** and selective data recording. For example, not every tick needs to be saved – maybe only major epochs or interesting events. Use binary serialization (the project added `bincode` for state saving ⁵²) for efficient storage of snapshots. Provide tools to analyze data on-the-fly to avoid storing raw data unnecessarily (for instance, compute summary statistics of the universe state each tick and log those, rather than keeping every particle's coordinates over time). This approach echoes practices in HPC simulations where typically one cannot save every time-step due to I/O limits – instead, one saves every Nth step and computes diagnostics during runtime.

- **Verification and Calibration vs Real Data:** An often overlooked aspect of fidelity is comparing simulation outcomes with known reality to ensure the models are tuned correctly. As EVOLVE incorporates all these physics, it should produce results that line up with known benchmarks: e.g., the cosmic microwave background (CMB) relic from recombination, the distribution of elements in the universe (around 75% H, 25% He after Big Bang nucleosynthesis), the stellar mass function, etc. While EVOLVE has broader goals (like evolving AI life), ensuring that up to the life emergence point the physics hasn't strayed from realism is important. This means calibrating parameters (like any subgrid models or stochastic models) to match observations or established simulations. For instance, if star formation is implemented via a simplified recipe, adjust it so that the simulated galaxies have plausible star formation rates. If life's emergence is too complex to simulate from pure chemistry, one might input at least the condition that life arises on Earth-like planets after X million years on average (based on Earth's fossil record) to calibrate the timeline. These are semi-empirical tweaks but will make the simulation more credible scientifically. Each such tweak should be grounded in literature (e.g., star formation recipes from astrophysical simulations, or biogenesis time estimates from astrobiology studies).

In summary, achieving maximal fidelity requires *focusing computational effort where it matters* and making careful approximations elsewhere. The project should continue to leverage **adaptive algorithms** and **parallel computing** to stretch the simulation's capabilities. It's a delicate balance: go too fine everywhere, and it never runs; go too coarse, and you miss important emergent phenomena. The solution is *graduated detail*, validation against known science, and a flexible architecture that can swap in more detailed models for subsystems as computing power grows in the future.

Best Practices for Building from First Principles

Building a universe simulation from first principles is akin to developing many simulations in one – a cosmology sim, a particle physics sim, a chemistry sim, an evolution sim, all integrated. Here are some **best practices** and methodological recommendations to manage this complexity:

- **Modularity and Interface Definition:** Keep the code modular by physics domains, as EVOLVE is doing. Clearly define interfaces between modules (e.g., how does the physics engine hand off information to the agent evolution engine?). Use abstraction barriers: the cosmology module doesn't need to know how the physics engine calculates a fusion reaction, only that, for example, a star of a given mass produces X energy. Conversely, the microphysics module can be oblivious to galactic dynamics. This encapsulation allows swapping out internals without breaking the whole system. For instance, if a more advanced physics model becomes available, it can replace the older one as long as it provides the same interface outputs (like reaction rates, force calculations, etc.). **Documentation** of these interfaces is crucial so developers (or open-source contributors) can work on separate pieces in parallel.
- **Incremental Validation:** Build confidence by testing each component on problems with known solutions. For physics: test the gravity integrator on a simple two-body orbit (compare to Kepler's laws), test the electromagnetic solver on known wave propagation or Coulomb's law in simple setups, test the nuclear network on known element production in stars (compare to standard stellar models). For the evolution aspect: test that digital organisms evolve in scenarios similar to known artificial life experiments (e.g., see if they exhibit open-ended innovation akin to work in Avida or other digital life platforms). Each module should have a suite of unit or integration tests anchored to

either analytical results or empirical data. This echoes a scientific approach: **verification and validation** steps ensure the simulation's building blocks are sound before trusting the full simulation.

- **Use of Units and Nondimensionalization:** Simulations can suffer from numerical precision issues when values vary by many orders of magnitude. The universe spans extremes (Planck energy $\sim 1.22 \times 10^{19}$ GeV versus cosmic microwave background photon energy $\sim 10^{-13}$ GeV, etc.). Using appropriate units or nondimensionalizing equations can help. For stability, one might internally use natural units ($c=1$, $\hbar=1$, $k_B=1$) in calculations, or scale variables to order 1 in code. Just be careful to convert back to meaningful units for output. Tracking units explicitly (possibly via a units library or simple wrappers) can prevent mistakes when mixing e.g. meters and centimeters, or seconds and years. It appears EVOLVE uses SI or astrophysical units consistently (with constants updated to CODATA values ⁴).
- **Randomness and Reproducibility:** Since many processes are stochastic (quantum events, mutations in evolution, etc.), the simulation should use a robust random number generator (RNG). It's wise to allow a **seed** to be set for reproducibility of a given run. Reproducibility is important for debugging and scientific analysis. If a run yields an interesting outcome (say, emergence of a certain life form), being able to replay it with the same seed ensures the phenomenon can be studied. Conversely, support running multiple simulations with different seeds (or even a different "random physics" if exploring multiverse ideas) to see distribution of outcomes.
- **Performance Profiling and Optimization:** Since performance is secondary but still relevant, regularly profile the simulation to find bottlenecks. This might reveal, for example, that the physics collision checking is eating most CPU, or memory usage spikes at certain times. Then focus optimization effort there – perhaps by using spatial indexing (like a grid or tree) for collisions instead of $O(n^2)$ checking, or by compressing data structures (Rust's ownership model can help avoid copies). The project has a diagnostics crate ² – that's excellent for monitoring performance in real-time. It could log metrics like loop iterations per second, memory footprint, etc., and perhaps trigger automatic adjustments (like reducing detail if memory is low – as hinted by a `--low-mem` mode in CLI ⁴⁶).
- **Community and Collaboration:** A project of this scope benefits from an open-source development model, inviting experts from different fields. The EVOLVE repository is public and already organized. As it grows, maintaining clear documentation (both user-facing and developer-facing) is key. A good practice is writing design docs for complex features (e.g. how the Planck-scale simulation is handled) and perhaps publishing them in the repository's `docs/` folder. This allows peer review of the approach even before coding is done, ensuring scientific rigor (peer-reviewed references should back those design decisions). Engaging with scientific communities (physicists, astrobiologists, AI researchers) via forums or conferences can provide feedback and new ideas. EVOLVE can also validate its components by comparing with established codes (for instance, cross-check its cosmology results with those from Gadget or Enzo on a standard test, or compare its agent evolution with known ALife experiments).
- **Ethical and Philosophical Considerations:** As a side note, simulating a universe including sentient AI agents raises philosophical questions (simulation ethics, etc.). While not the focus of the technical plan, it's worth being mindful of this especially if the simulation is intended to potentially produce

conscious entities (the documentation even lists “emergence of consciousness” as a milestone ²⁰). In practical terms, this means providing controls to pause or stop the simulation, introspect agent states, and ensure that researchers can analyze or intervene in potentially unsafe AI behaviors. This parallels discussions in AI safety and the simulation hypothesis. By building in an “**observer mode**” or constraints (like not letting simulated super-intelligences escape their sandbox, as fanciful as that sounds), the project stays within bounds of research and avoids unintended consequences.

Following these best practices will guide EVOLVE to evolve (pun intended) into a robust platform that not only pushes the boundaries of simulation fidelity but does so in a controlled, validated, and transparent manner.

Relevant Open-Source Simulation Frameworks (Comparison Table)

To leverage existing knowledge and possibly code, here is a table of notable open-source simulation libraries and engines relevant to EVOLVE’s goals. These can serve either as **integration targets** (linking or borrowing code) or simply as **references** for methodologies and validation:

Name & Source	Domain/ Physics	Description & Use in Context	Trade-offs
Einstein Toolkit ³⁸	General Relativity (Astrophysics)	C++/Fortran framework for simulating relativistic spacetimes (black holes, neutron stars). Provides solvers for Einstein’s field equations and hydrodynamics under strong gravity. Useful for incorporating realistic gravity and gravitational wave dynamics into specific scenarios (e.g., mergers).	High complexity and computational cost; uses grid methods requiring HPC. Would need coupling to rest of sim (not turnkey for whole universe).
GADGET-2/4 ³⁶	Cosmology, N-Body Gravity + SPH	Massively parallel N-body code with smoothed-particle hydrodynamics. Used for large-scale structure and galaxy formation simulations ³⁶ . Could inspire EVOLVE’s handling of gravity and gas dynamics on cosmic scales.	Gravity+hydro only (no quantum/chemistry or feedback without additional modules). Written in C; EVOLVE would need to reimplement or interface via data exchange.

Name & Source	Domain/ Physics	Description & Use in Context	Trade-offs
RAMSES ⁴⁸	Astrophysical Fluid Dynamics (AMR)	An open-source Eulerian code with Adaptive Mesh Refinement for self-gravitating magnetized flows. Good for high-resolution zoom-in simulations (e.g., star formation regions in a galaxy).	Grid-based – can handle fluids & fields well, but particle support for e.g. star particles needed. Could be too slow if whole universe on a fine grid.
Geant4 ³⁹	Particle Physics (High Energy & Nuclear)	Toolkit for simulating particles through matter ³⁹ . Includes comprehensive physics lists for electromagnetic, weak, strong interactions from eV to TeV scales. In EVOLVE, it could be used to handle radiation transport, cosmic ray interactions, and detailed particle collision outcomes within environments.	Very fine-grained and C++ based; integrating it for every particle would be overkill. Better used selectively (e.g., for a high-energy experiment within the sim, or precomputed lookup tables). Also, steep learning curve.
LAMMPS ⁵³	Molecular Dynamics (Classical)	Widely used open-source molecular dynamics code for materials science ⁵⁴ . Supports various interatomic potentials and runs on parallel hardware. Could simulate chemistry or material behavior (e.g., protocell membranes, or crystalline planet cores) at atomistic or coarse-grained level within EVOLVE.	Classical MD only (no electronic structure). Not feasible to embed for large atom counts due to speed. Possibly use for small sub-simulations (like testing a biochemical reaction in detail then feeding results to EVOLVE).
OpenFOAM	Continuum Fluid Dynamics (CFD)	Open-source fluid solver for gases and liquids (Navier-Stokes equations), with plugins for heat transfer, chemical reactions, etc. Could be employed for high-fidelity simulation of, say, a planetary atmosphere or ocean circulation on an Earth-analog planet in EVOLVE.	CFD at global scale is heavy; would need coarse grids. Not inherently coupled with other physics like gravity (though that can be added). Possibly overkill unless focusing on climate or fluid dynamics specifically.

Name & Source	Domain/ Physics	Description & Use in Context	Trade-offs
Bullet Physics	Classical Rigid- Body Dynamics (Game Engine)	Real-time physics engine (C++) commonly used in games for collision and rigid body simulation. Could handle macroscopic structures (spacecraft, debris, etc.) efficiently and with approximate accuracy.	Not high-precision (tuned for speed and stability over accuracy). Lacks units (user must scale to avoid floating-point issues). Not suitable for microphysics or relativity. Use only if simulating mechanical systems at human scale with performance needs.
Avida / Open Worm etc.	Artificial Life & Biologic Simulation	Avida is a digital evolution platform where self-replicating programs mutate and evolve (useful for studying open-ended evolution) – it's not a physics sim, but could inform EVOLVE's agent evolution logic. OpenWorm simulates a real organism's cells and neural network. These represent biological layers that EVOLVE might emulate when bridging from physics to life.	These are domain-specific: Avida ignores physics (digital grid world), OpenWorm is limited to one worm's biology. Rather than integrate, EVOLVE can draw algorithms (e.g., genetic mutation routines, metabolic network simulation) from them.
Illustris / Arepo <small>
 55</small>	Galaxy Formation (Multiphysics)	The Illustris project's code (Arepo) is a moving-mesh simulation including gravity, hydrodynamics, star formation, and feedback processes (supernova, AGN feedback). It's not fully open-source, but methods are published. Arepo's approach of a dynamic mesh combining Lagrangian and Eulerian benefits might inspire EVOLVE's handling of fluids and space adaptivity.	Arepo is complex and not readily available. It excels at specific astrophysical realism (e.g., realistic galaxy morphologies), which EVOLVE could strive for if focusing on that era. But Illustris still doesn't include quantum microphysics or life – its scope is narrower (cosmic structure to perhaps simple chemical enrichment).

Name & Source	Domain/ Physics	Description & Use in Context	Trade-offs
MESA (Modules for Experiments in Stellar Astrophysics)	Stellar Evolution (1D)	An open-source 1D stellar evolution code that models stars from formation to supernova by solving thermodynamics, nuclear networks, and convection in shells. Could be used to validate EVOLVE's star simulations or even integrated to provide realistic star life-cycle data (mass–luminosity relationships, etc.) ⁹ .	1D approximation (assumes spherical symmetry), so not spatially resolved. But very detailed in microphysics of stars. Might be too specialized to integrate fully, but can calibrate EVOLVE's simpler star models.

Notes: This table isn't exhaustive, but highlights the diversity of tools. EVOLVE's challenge is essentially to unify aspects of all these domains. While direct code integration might sometimes be infeasible (due to language differences or performance), studying these projects' methodologies is invaluable. For example, the **Einstein Toolkit** can guide handling of relativistic coordinates and gauge conditions; **GADGET** provides algorithms for long-range gravity and SPH hydrodynamics with excellent efficiency; **Geant4** offers a repository of physics knowledge for microscale events; **LAMMPS** and others provide models for materials and molecular interactions, etc. Adopting proven algorithms from them (or even data files, like opacities or cross-sections) will accelerate development and ensure scientific correctness.

Conclusions

Creating a simulation of the universe with **perfect fidelity** is a herculean task – one that pushes the limits of physics, computation, and even philosophy. The EVOLVE project has made commendable progress in outlining and beginning to implement this vision: a modular architecture capturing everything from quantum particle interactions to the rise of intelligent life ¹⁵ ³⁴.

Moving forward, the development should continue to be **grounded in peer-reviewed science**, using known physics wherever possible and introducing speculative elements cautiously and transparently. We have highlighted how all the essential physics can be incorporated: general relativity for gravity, quantum mechanics and the Standard Model for microphysics, thermodynamics to bridge scales, and even potential new physics as optional modules. The use of existing open-source engines and scientific codes – either through inspiration or integration – will greatly benefit EVOLVE, preventing the need to reinvent decades of research.

Key recommendations include employing adaptive multi-scale techniques to handle the vast range of scales, validating each component against known benchmarks, and embracing performance optimizations like parallelism and level-of-detail adjustments to make the simulation tractable on a high-end server. It's also prudent to keep the simulation flexible: users might run a “universe” at reduced scale or accelerated time, for example, to observe particular phenomena (like a miniature laboratory universe), so the design should allow configurability (which the CLI's presets hint at ⁴⁶).

Ultimately, beyond the technicalities, maintaining the **scientific spirit** is vital. The simulation should be used as a tool to experiment with hypotheses about the universe: How do galaxies form? How might life originate under different conditions? What if fundamental constants were different? By faithfully implementing the laws of nature as we know them (and offering extensions for those we speculate), EVOLVE can become not just a software project but a **virtual laboratory** for cosmology, physics, and evolution. It could enable unprecedented experiments – for instance, rewinding and altering the course of cosmic history (“**rewind**” and “**time-warp**” functions are already planned ⁵⁶ ²⁷) – to see what outcomes result, a feat impossible in the real world.

In conclusion, building the most advanced universe simulation is an iterative, multi-disciplinary effort. By heeding the methodologies and resources outlined here, the EVOLVE team can methodically expand the simulation’s fidelity. Each added layer of physics or complexity will bring it a step closer to the holy grail: a simulation in which, given fundamental laws and initial conditions, **the cosmos unfolds naturally**, birthing stars, galaxies, and potentially life and intelligence, all within a self-consistent digital universe. This pursuit not only tests our understanding of physics across scales but also illuminates the path to creating **true artificial life** grounded in fundamental reality – a profound scientific frontier.

Sources: The recommendations and comparisons above draw from established scientific literature and existing simulation frameworks. For instance, the computational limits of a universe simulation have been discussed by Lloyd (2002), who estimated $\sim 10^{120}$ operations and 10^{90} bits as a lower bound for a full simulation ⁴⁴ ⁴⁵. Modern cosmological codes like GADGET ³⁶ and adaptive mesh refiners ⁵⁷ ⁴⁸ demonstrate techniques for handling gravity and fluids at scale, while particle physics toolkits like Geant4 encapsulate the Standard Model interactions over a wide energy range ³⁹. Quantum gravity approaches (CDT, LQG) suggest spacetime’s behavior near the Planck scale, e.g. emergent lower dimensions ⁴³, which can inform how we discretize space-time at small scales. The EVOLVE project’s own documentation and commit history have been referenced to align these recommendations with the project’s current trajectory ⁴ ⁹ ¹⁶. By synthesizing insights from these sources, the path forward for EVOLVE becomes clearer and firmly rooted in scientific consensus and state-of-the-art simulation practices.

1 2 3 4 12 13 14 15 17 18 19 20 21 24 25 26 27 30 31 32 34 35 42 46 49 50 51 56 GitHub

- expressionneuroscienceinstitute/evolve at dev

<https://github.com/expressionneuroscienceinstitute/evolve/tree/dev>

5 6 7 8 9 10 11 16 22 23 28 29 33 52 Commits · expressionneuroscienceinstitute/evolve · GitHub

<https://github.com/expressionneuroscienceinstitute/evolve/commits/dev/>

36 38 41 47 48 57 GitHub - pmocz/awesome-astrophysical-simulation-codes: A curated list of awesome astrophysical simulation codes

<https://github.com/pmocz/awesome-astrophysical-simulation-codes>

37 List of Astrophysics Codes

<https://results-preview-punch4nfdi.sirrah.aip.de/?md=/docs/Compute/Codes/astro-codes.md>

39 GEANT4 the physics simulation toolkit | symmetry magazine

https://www.symmetrymagazine.org/article/november-2005/geant4-physics-simulation-toolkit?language_content_entity=und

40 Prebiotic chemistry and origins of life research with atomistic ...

<https://www.sciencedirect.com/science/article/pii/S157106451830109X>

43 Causal dynamical triangulation - Wikipedia

https://en.wikipedia.org/wiki/Causal_dynamical_triangulation

44 45 THE COMPUTATIONAL UNIVERSE | Edge.org

https://www.edge.org/conversation/seth_lloyd-the-computational-universe

53 54 LAMMPS: Open-Source, High-Performance, and High ...

<https://www.energy.gov/eere/h2awsm/lammps-open-source-high-performance-and-high-fidelity-molecular-dynamics-code>

55 Astrophysics Codes - Forschungszentrum Jülich

<https://www.fz-juelich.de/en/ias/jsc/about-us/structure/simulation-and-data-labs/sdl-physics-and-astronomy/simulation-codes>